

## ✓ Lecture 12 - Loops - Apply Family

---

### ✓ Packages

```
# none
```

---

### ✓ Apply family of functions

- The prior lecture covered loops, which is a mechanism used in many other languages
  - R conveniently has a family of functions called the `apply` family, which performs implicit looping
  - Most of the time, we can use these functions instead of creating our own for loops!
- 
- Today, we will cover the following four apply family functions
    - `apply()`
      - It takes a data structure, a margin (specifying whether to operate along rows or columns), and a function to apply. It returns a vector or matrix depending on the input and the margin.
    - `lapply()` - List Apply
      - This function applies a function to each element of a list, returning a list of the results. It maintains the structure of the input list.
    - `sapply()` - Simple Apply
      - Similar to `lapply`, but it attempts to simplify the output if possible, returning a vector or matrix instead of a list. This is useful when the output of the function is consistent across all elements.

- `tapply` - Table Apply
  - This function applies a function to elements grouped by a factor or categorical variable, returning a table of results.
- We will also cover the powerful `do.call()` function

## ▼ `apply()`

- The `apply()` function applies a function to the rows or the columns of a matrix or dataframe and outputs a vector or list

### apply: Apply Functions Over Array Margins

#### Description

Returns a vector or array or list of values obtained by applying a function to margins of an array or matrix.

#### Usage

```
apply(X, MARGIN, FUN, ...)
```

#### Arguments

<code>X</code>	an array, including a matrix.
<code>MARGIN</code>	a vector giving the subscripts which the function will be applied over. E.g., for a matrix <code>"1"</code> indicates rows, <code>"2"</code> indicates columns, <code>"c(1, 2)"</code> indicates rows and columns. Where <code>"X"</code> has named dimnames, it can be a character vector selecting dimension names.
<code>FUN</code>	the function to be applied: see 'Details'. In the case of functions like <code>"+"</code> , <code>"%*%"</code> , etc., the function name must be backquoted or quoted.
<code>...</code>	optional arguments to <code>"FUN"</code> .

- Find the mean of each row of the `mtcars` data frame

```
# view first few lines of data frame
head(mtcars)
```



A data.frame: 6 × 11

```
      mpg    cyl  disp    hp  drat    wt    qsec    vs    am    gear    (
```

	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
<b>Mazda RX4</b>	21.0	6	160	110	3.90	2.620	16.46	0	1	4
<b>Mazda RX4 Wag</b>	21.0	6	160	110	3.90	2.875	17.02	0	1	4
<b>Datsun 710</b>	22.8	4	108	93	3.85	2.320	18.61	1	1	4
<b>Hornet 4 Drive</b>	21.4	6	258	110	3.08	3.215	19.44	1	0	3
<b>Hornet</b>	18.7	8	360	175	3.15	3.440	17.02	0	0	0

```
# find mean of each row, 1 = row
rowmeans <- apply(mtcars, 1, mean)
print(rowmeans)
```

Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
29.90727	29.98136	23.59818	38.73955
Hornet Sportabout	Valiant	Duster 360	Merc 240D
53.66455	35.04909	59.72000	24.63455
Merc 230	Merc 280	Merc 280C	Merc 450SE
27.23364	31.86000	31.78727	46.43091
Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
46.50000	46.35000	66.23273	66.05855
Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla
65.97227	19.44091	17.74227	18.81409
Toyota Corona	Dodge Challenger	AMC Javelin	Camaro Z28
24.88864	47.24091	46.00773	58.75273
Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
57.37955	18.92864	24.77909	24.88027
Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E
60.97182	34.50818	63.15545	26.26273

Use `data.frame()` to convert the output vector to a dataframe.

```
# find mean of each row
rowmeans <- apply(mtcars, 1, mean)
rowmeans <- data.frame(rowmeans)
head(rowmeans)
```

```
A data.frame: 6 × 1
      rowmeans
      <dbl>
1 Mazda RX4 29.90727
2 Mazda RX4 Wag 29.98136
```

<b>Mazda RX7 Wag</b>	23.59180
<b>Datsun 710</b>	23.59818
<b>Hornet 4 Drive</b>	38.73955
<b>Hornet Sportabout</b>	53.66455
<b>Valiant</b>	35.04909

- Note that the above doesn't make much sense!
- It makes more sense to average each column of the data frame since values within a column are the same units of measurement

```
# find mean of each column, 2 = column
colmeans <- apply(mtcars, 2, mean)
print(colmeans)
```

mpg	cyl	disp	hp	drat	wt	qsec
20.090625	6.187500	230.721875	146.687500	3.596563	3.217250	17.848750
vs	am	gear	carb			
0.437500	0.406250	3.687500	2.812500			

Use `data.frame()` to convert the output vector to a dataframe.

```
# find mean of each column
colmeans <- apply(mtcars, 2, mean)
colmeans <- data.frame(colmeans)
```

```
head(colmeans)
```

A data.frame: 6 × 1

	<b>colmeans</b>
	<b>&lt;dbl&gt;</b>
<b>mpg</b>	20.090625
<b>cyl</b>	6.187500
<b>disp</b>	230.721875
<b>hp</b>	146.687500
<b>drat</b>	3.596563
<b>wt</b>	3.217250

## ▼ lapply()

- The `lapply()` function applies a function to each element of a list and outputs another list

### lapply: Apply a Function over a List or Vector

#### Description

``lapply`` returns a list of the same length as ``x``, each element of which is the result of applying ``FUN`` to the corresponding element of ``x``.

``sapply`` is a user-friendly version and wrapper of ``lapply`` by default returning a vector, matrix or, if ``simplify = "array"``, an array if appropriate, by applying ``simplify2array()``.

``sapply(x, f, simplify = FALSE, USE.NAMES = FALSE)`` is the same as ``lapply(x, f)``.

``vapply`` is similar to ``sapply``, but has a pre-specified type of return value, so it can be safer (and sometimes faster) to use.

``replicate`` is a wrapper for the common use of ``sapply`` for repeated evaluation of an expression (which will usually involve random number generation).

``simplify2array()`` is the utility called from ``sapply()`` when ``simplify`` is not false and is similarly called from ``mapply()``.

#### Usage

```
lapply(X, FUN, ...)  
sapply(X, FUN, ..., simplify = TRUE, USE.NAMES = TRUE)  
vapply(X, FUN, FUN.VALUE, ..., USE.NAMES = TRUE)  
replicate(n, expr, simplify = "array")  
simplify2array(x, higher = TRUE)
```

#### Arguments

**X** a vector (atomic or list) or an ``expression`` object. Other objects (including classed objects) will be coerced by ``base::as.list``.

**FUN** the function to be applied to each element of ``x``: see 'Details'. In the case of functions like ``+``, ``%%``, the function name must be backquoted or quoted.

- Remember that a dataframe is a structured list!
- We can use `lapply()` find the mean of each variable in the data frame

```
# Find th mean of each column  
lapply(mtcars, mean)
```

```
$mpg  
20.090625
```

```

$cyl
6.1875
$disp
230.721875
$hp
146.6875
$drat
3.5965625
$wt
3.21725
$qsec
17.84875
$vs
0.4375
$am
0.40625
$gear
3.6875
$carb
2.8125

```

&gt;

```

# Converting unstructured list to a data frame
data.frame(lapply(mtcars, mean))

```

A data.frame: 1 × 11

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	g
<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<d
20.09062	6.1875	230.7219	146.6875	3.596563	3.21725	17.84875	0.4375	0.40625	3.6

```
sapply()
```

```
[ ] ↪ 7 cellules masquées
```

▼ `tapply()`

- The `tapply()` functions applies a function for each factor level in a variable
- This function is very useful!

## tapply: Apply a Function Over a Ragged Array

**Description**

Apply a function to each cell of a ragged array, that is to each (non-empty) group of values given by a unique combination of the levels of certain factors.

**Usage**

```
tapply(X, INDEX, FUN = NULL, ..., default = NA, simplify = TRUE)
```

**Arguments**

<b>X</b>	an R object for which a <code>`split`</code> method exists. Typically vector-like, allowing subsetting with <code>`[`</code> .
<b>INDEX</b>	a <code>`list`</code> of one or more <code>`factor`</code> s, each of same length as <code>`X`</code> . The elements are coerced to factors by <code>`as.factor`</code> .
<b>FUN</b>	a function (or name of a function) to be applied, or <code>`NULL`</code> . In the case of functions like <code>`+`</code> , <code>`%*%`</code> , etc., the function name must be backquoted or quoted. If <code>`FUN`</code> is <code>`NULL`</code> , <code>tapply</code> returns a vector which can be used to subscript the multi-way array <code>`tapply`</code> normally produces.

- For example, we can use `tapply()` to calculate the average mpg for each cylinder engine

```
# means for each cylinder
cyl_means <- tapply(mtcars$mpg, mtcars$cyl, mean)
print(cyl_means)
```

```
      4      6      8
26.66364 19.74286 15.10000
```

## ✓ Custom Functions and the `apply()` Family

- Thus far, all we have done is calculate averages
- However, the `apply()` family can use different functions as well...even custom functions!
- The major challenge is understanding
  1. the appropriate inputs to the function
  2. the expected output of the function

## ▼ Examples

- Finding the maximum of each variable using `apply()`

```
# find maximum of each variable
print(apply(mtcars, 2, max))
```

```
      mpg      cyl    disp      hp    drat      wt     qsec      vs      am     gear
33.900    8.000  472.000  335.000   4.930   5.424   22.900    1.000    1.000    5.000
 carb
 8.000
```

- Finding the minimum of each variable using `lapply()`

```
lapply(mtcars, max)
```

```
$mpg
33.9
$cyl
8
$disp
472
$hp
335
$drat
4.93
$wt
5.424
$qsec
22.9
$vs
1
$am
1
$gear
5
$carb
8
```

- Finding the median of each variable using `sapply()`

```
print(sapply(mtcars, max))
```



mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear
33.900	8.000	472.000	335.000	4.930	5.424	22.900	1.000	1.000	5.000
carb									
8.000									

- Finding the quantiles of the `mpg` variable by the number of cylinders using `tapply()`

```
print(tapply(mtcars$mpg, mtcars$cyl, quantile))
```

```
$`4`
 0%  25%  50%  75% 100%
21.4 22.8 26.0 30.4 33.9

$`6`
 0%  25%  50%  75% 100%
17.80 18.65 19.70 21.00 21.40

$`8`
 0%  25%  50%  75% 100%
10.40 14.40 15.20 16.25 19.20
```

## ▼ Custom Function Example

- We can build our own function to summarize the data
- For example, `mean()`, `min()`, and `max()` are built-in functions in R, but we can build our own functions to be used by the `apply` family

```
# define our own function
vector_summary <- function(x) {

  # calculate descriptive statistics
  mean_out <- mean(x)
  min_out  <- min(x)
  max_out  <- max(x)

  # combine into a vector
  output <- c(mean_out, min_out, max_out)

  # return summary
```

```

    # return summary
    return(output)
}

# apply function to mpg for each unique cylinder
mpg_summary <- tapply(mtcars$mpg, mtcars$cyl, vector_summary)
mpg_summary

```

```

$`4`
26.6636363636364 · 21.4 · 33.9
$`6`
19.7428571428571 · 17.8 · 21.4
$`8`
15.1 · 10.4 · 19.2

```

## ▼ do.call()

- The function `do.call()` can be used to combine elements of a list using a specified function
- For example, if we want to concatenate the elements of a list

## do.call: Execute a Function Call

### Description

`do.call` constructs and executes a function call from a name or a function and a list of arguments to be passed to it.

### Usage

```
do.call(what, args, quote = FALSE, envir = parent.frame())
```

### Arguments

<code>what</code>	either a function or a non-empty character string naming the function to be called.
<code>args</code>	a <i>list</i> of arguments to the function call. The <code>names</code> attribute of <code>args</code> gives the argument names.

- For example, we can use `do.call()` to combine the following output into a data summary table

```
# summary table
mpg_summary <- tapply(mtcars$mpg, mtcars$cyl, vector_summary)
mpg_summary
```

```
$`4`
26.6636363636364 · 21.4 · 33.9
$`6`
19.7428571428571 · 17.8 · 21.4
$`8`
15.1 · 10.4 · 19.2
```

```
# row bind each summary in the unstructured list
mpg_summary <- do.call(rbind, mpg_summary)
mpg_summary
```

```
A matrix: 3 × 3 of type dbl
 4 26.66364 21.4 33.9
 6 19.74286 17.8 21.4
 8 15.10000 10.4 19.2
```

- We can now continue organizing the output for viewing

```
# convert to a data frame
mpg_summary <- data.frame(mpg_summary)
mpg_summary
```

```
A data.frame: 3 × 3
      X1      X2      X3
  <dbl> <dbl> <dbl>
1 26.66364 21.4 33.9
2 19.74286 17.8 21.4
3 15.10000 10.4 19.2
```

```
# change column names
names(mpg_summary) <- c("mean", "min", "max")
```

```
mpg_summary
```

```
A data.frame: 3 × 3
```

	mean	min	max
	<dbl>	<dbl>	<dbl>
<b>4</b>	26.66364	21.4	33.9
<b>6</b>	19.74286	17.8	21.4
<b>8</b>	15.10000	10.4	19.2

```
# change row names
```

```
rownames(mpg_summary) <- paste(c(4, 6, 8), "cyl", sep="")
```

```
mpg_summary
```

```
A data.frame: 3 × 3
```

	mean	min	max
	<dbl>	<dbl>	<dbl>
<b>4cyl</b>	26.66364	21.4	33.9
<b>6cyl</b>	19.74286	17.8	21.4
<b>8cyl</b>	15.10000	10.4	19.2